

OJXPERF: Pinpointing Object-level Redundancies in Java Programs

Anonymous Author(s)

ABSTRACT

This document provides the supplementary materials for the PLDI 2020 submission. It shows the overall overhead analysis. Then this document describes the optimizations under the guidance of OJXPERF for JGFSerailBench [2] and RoaringBitmap [1].

1 CASE STUDIES

1.1 JGFSerailBench

JGFSerailBench is a benchmark of Java Grande benchmark suite, which are designed to be representative of scientific and other numerically intensive computation [2]. We run JGFSerailBench using its default input size. After profiling JGFSerailBench, one object, arraybase, which is allocated at line 371 in method JGFrUn of class JGFSerailBench stood out, as shown in Listing 1. Listing 1 shows the allocation site of this arraybase is in a while loop. In each loop, the program creates a new arraybase and puts some contents into this arraybase. As this arraybase is a two-dimensional array, arraybase uses writeObject method (line 375) to put its one-dimensional elements (can be seen as an object) into an object arrayout. OJXPERF reports that the redundancy factor θ of the created arraybase is 68.8%, and its lower bound factor ω indicates that the largest identical arraybase group size is larger than 15.2%. To avoid such redundant object copy operations, we redesign the code as we first check whether the program has a different arraybase compared with the one from the previous iteration. If the arraybase keeps the same, then copy of arraybase's contents is unnecessary. This optimization yields a $(1.1 \pm 0.03)\times$ speedup to the entire program.

```
368 public void JGFrUn(){
369     ...
370     while (time < TARGETTIME && size < MAXSIZE){
371         ▶arraybase = new arrayitem [size][LENGTH];
372         arrayout = new ObjectOutputStream(arrayout);
373         ... //put contents in array base
374         for (j=0; j<size; j++) {
375             ▶arrayout.writeObject( arraybase[j]);
376         }
377     }
378     ...
379 }
```

Listing 1: OJXPERF pinpoints the arraybase array suffering from object-level redundancy in JGFSerailBench

```
284 public void deserialize(DataInput in) {
285     ...
286     final short keys[] = new short[this.size];
287     for (int k = 0; k < this.size; ++k) {
288         keys[k] = Short.reverseBytes(in.readShort());
289         ...
290         ▶this.keys[k] = keys[k];
291     }
292     ...
293 }
```

Listing 2: The source code highlighted by OJXPERF shows the keys array suffering from object-level redundancy in RoaringBitmap.

1.2 RoaringBitmap

Roaring bitmap are compressed bitmaps which tend to outperform conventional compressed bitmaps such as WAH, EWAH or Concise [1]. We run RoaringBitmap using SerializationBenchmark as input. OJXPERF reports an object, array keys, which is accessed at line 290 in method deserialize of class MutableRoaringArray, as shown in Listing 2. In this deserialize method, the program declares a short array keys (line 286) and fills every slots of this array (line 288). For each iteration, after the program obtained the value of keys[k], it will update the class member array keys[k]'s value at end of the loop. OJXPERF reports that the array keys always equals with the array keys from the last iteration (redundancy factor θ is 100%), which means that the method deserialize often generates the exactly same array keys comparing with the one from the last call of this method. We redesign the code as instead of updating the array keys value at end of every loop, the program only updates the array keys outside this loop using the Arrays.copyOf method if the deserialize method generates a different array keys. The score's unit used in RoaringBitmap is microseconds per operation (throughput). While this simple fix did not lead to a significant running time reduction (only 2.1%), the optimization increases the throughput by $(1.09 \pm 0.01)\times$.

REFERENCES

- [1] RoaringBitmap authors. 2019. RoaringBitmap. <https://github.com/RoaringBitmap/RoaringBitmap>.
- [2] Mark Bull, Lorna Smith, Martin Westhead, David Henty, and Robert Davey. 2001. Java Grande benchmark suite. <https://www.epcc.ed.ac.uk/research/computing/performance-characterisation-and-benchmarking/java-grande-benchmark-suite>.